Sruthijha Pagolu

+1 (850) 508-7178 | sruthipagolu@gmail.com | linkedin.com/in/sruthijha | github.com/Sruthijha

Education

University of Colorado Boulder — Boulder, CO
Master of Science in Data Science; ***GPA: 3.96/4.0*** — Aug 2023 – Dec 2025

Mahindra University — Hyderabad, India
Bachelor of Technology in Electrical and Electronics Engineering; ***GPA: 3.54/4.0*** — Jul 2019 – May 2023

Technical Skills

Languages: Python (Pandas, NumPy, PySpark), SQL, R, Git, FastAPI, SQLAlchemy
Data & ML: scikit-learn, statistics, A/B testing, Faster-Whisper, LangChain, LangGraph, RAG, prompt engineering
Platforms/Cloud: AWS (Lambda, S3, SQS, DynamoDB, API Gateway), Azure (Data Factory, Databricks, Synapse), Snowflake, Docker, Terraform, pytest, GitHub Actions, n8n
Visualization: Power BI, Tableau, Looker, Streamlit

Experience

Breaking Games — Remote
Data Analytics Extern — May 2026 – Jul 2026

- Consolidated six unreconciled data sources into a queryable database and mapped a data architecture tracing the customer journey from ad click to purchase, supporting a company managing $37K in Meta ad spend against $61K in H1 sales with new-customer growth down 48% YoY.
- Assembled a product performance leaderboard from Shopify sales data using SQL aggregation (revenue, units sold, AOV, RANK() OVER for revenue rank) to support leadership's decision on which games to reprint ahead of the holiday season, verifying aggregate accuracy through independent cross-checks before reporting figures.

Goodie Bag Food Co. — New York City, New York
Data Intern — Jan 2026 – Apr 2026

- Modeled a weighted Acquisition Potential Score across 86 NYC neighborhoods (renter occupancy, college population, subway density) using Census and NTA data, catching and removing an implausible signal during a data quality audit before it skewed the model.
- Layered competitive saturation and independent-shop supply onto the scoring model to assign each neighborhood a launch-readiness status, then ranked 2,079 shops across 36 priority neighborhoods to produce a partner-outreach target list culminating in a Brooklyn-first launch recommendation delivered to the COO.

Intech Insurance Surveyors — Hyderabad, India
Data Scientist — Jul 2022 – Jun 2023

- Engineered a Python and SQL pipeline to clean, standardize, and consolidate 5,000+ multi-source insurance claim records, then generated features (claim amount deviation, claim frequency, policy duration) for predictive modeling.
- Trained and evaluated Logistic Regression and ensemble Random Forest models with SMOTE for class imbalance, assigning risk scores that helped surveyors prioritize claims for investigation.
- Delivered interactive Power BI dashboards covering branch-wise claim volume, severity, and monthly loss trends, replacing manual Excel reporting and giving management a consolidated operational view for decision-making.

Projects

Research Paper Assistant: RAG-Based Document Q&A System — Sep 2025 – Dec 2025

- Architected a Retrieval-Augmented Generation pipeline over unstructured document text using LangChain and OpenAI embeddings, returning grounded answers with source citations; tuned chunking and retrieval strategies to balance accuracy and latency.
- Assembled an iterative evaluation framework to measure retrieval reliability and answer accuracy across pipeline versions, incorporating bias checks to validate output quality.

Minutes Maker: AI-Powered Meeting Automation Pipeline — Jan 2025 – Apr 2025

- Designed and deployed a serverless ETL pipeline (AWS Lambda, S3, SQS, DynamoDB) that transcribes meeting recordings via Faster-Whisper and generates structured minutes with an OpenAI-based summarizer.
- Containerized the pipeline with Docker, managed infrastructure via Terraform, and ran automated testing and CI/CD with pytest and GitHub Actions; live demo at minutes-maker.streamlit.app.

Real-Time Emotion Recognition from Speech — Jan 2024 – May 2024

- Trained a hybrid CNN and LSTM model to classify eight emotional states from speech using the RAVDESS dataset, extracting MFCCs, chroma features, and mel-spectrograms to capture spectral and temporal patterns.
- Applied data augmentation (noise injection, time stretching, pitch shifting) and tuned training with the Adam optimizer, achieving 95% classification accuracy across emotion categories.